



UNIMORE

UNIVERSITÀ DEGLI STUDI DI  
MODENA E REGGIO EMILIA


# Leader Election Algorithms

---

# The Bully Algorithm - Overview

---

- Process  $p$  calls an election when it notices that the coordinator is no longer responding.
- High-numbered processes “bully” low-numbered processes out of the election, until only one process remains.
- When?
  - When a crashed process reboots, it holds an election. If it is now the highest-numbered live process, it will win.
  - When some process recognize that the current leader is no longer active, I calls an election






Process  $p$  sends an election message to all *higher-numbered* processes in the system. If no process responds, then  $p$  becomes the coordinator. If a higher-level process ( $q$ ) responds, it sends  $p$  a OK message that terminates  $p$ 's role in the algorithm

The process  $q$  now calls an election (if it has not already done so).

Repeat until no higher-level process responds. The last process to call an election “wins” the election.

The winner sends a message to other processes announcing itself as the new coordinator.



If 7 comes back on line, it will call an election

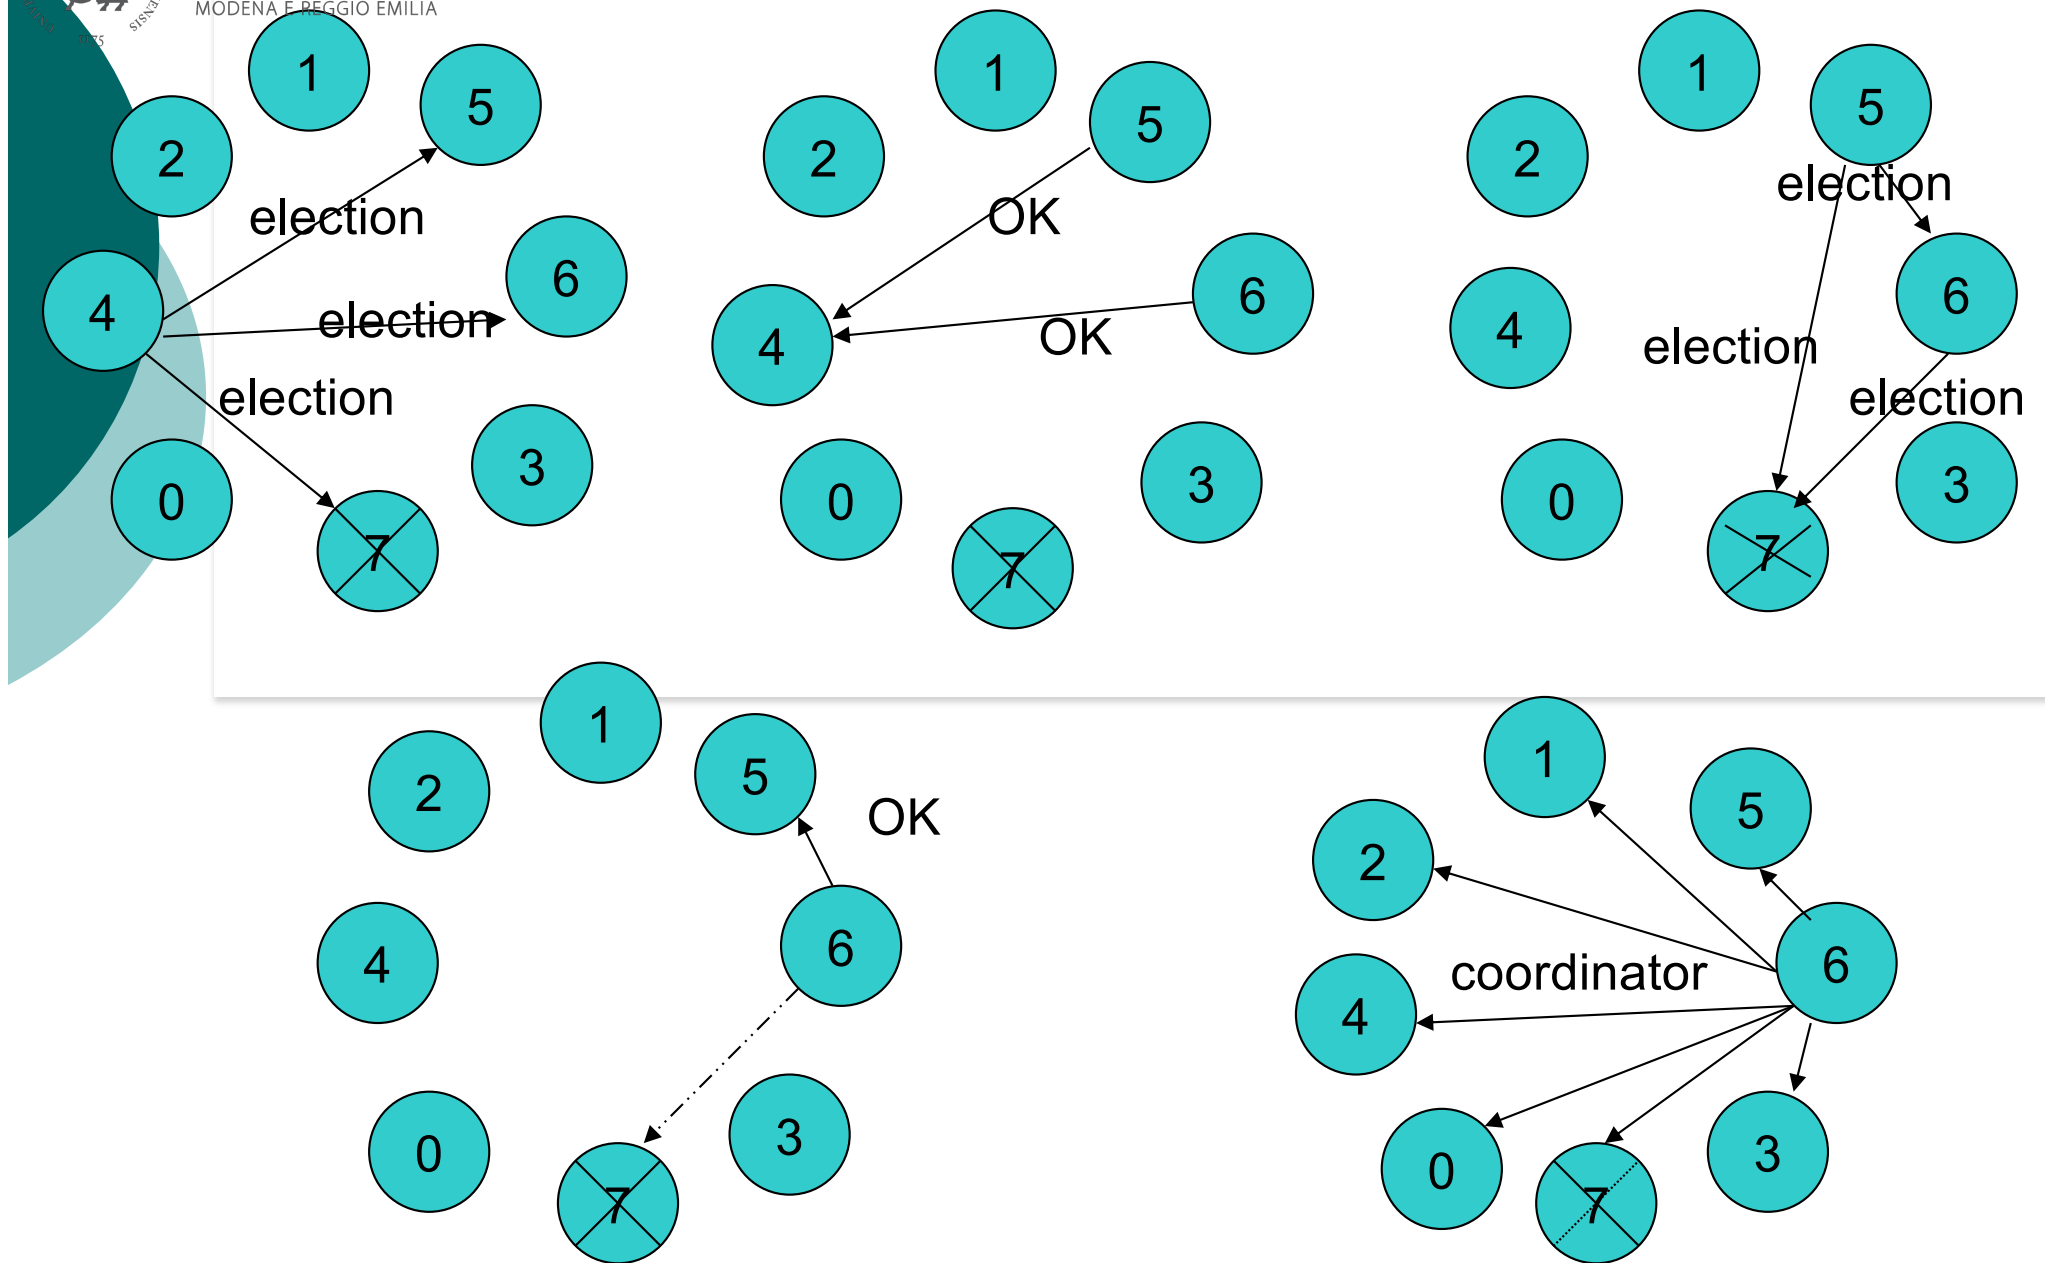


Figure 6-20

# Analysis

---

- Works best if communication in the system has bounded latency so processes can determine that a process has failed by knowing the upper bound (UB) on message transmission time (T) and message processing time (M).
  - $UB = 2 * T + M$
- However, if a process calls an election when the coordinator is still active, the coordinator will win the election.



## Complexity

---

- The complexity (i.e., the number of messages required to complete) depends on who starts the election
- Given we have  **$n$**  processes
- If the election is started by the highest ID process
  - **$n-1$**  messages are enough (“I will be the leader” said to everyone else!!)
- If the election is started by the lowest ID process
  - **$n(n-1)$**  messages

# The Ring Algorithm - Overview


---

- The ring algorithm assumes that the processes are arranged in a logical ring and each process knows the order of the ring of processes.
- Processes are able to “skip” faulty systems: instead of sending to process  $j$ , send to  $j + 1$ .
- Faulty systems are those that don't respond in a fixed amount of time.



# A Ring Algorithm

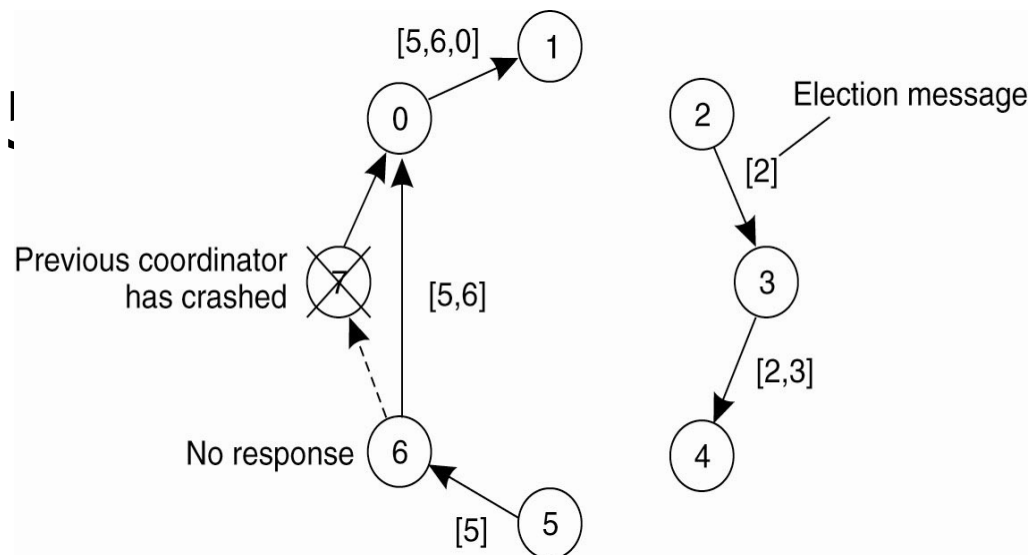
- P thinks the coordinator has crashed; builds an ELECTION message which contains its own ID number.
- Sends to first live successor
- Each process adds its own number and forwards to next.
- OK to have two elections at once.



- When the message returns to  $p$ , it sees its own process ID in the list and knows that the circuit is complete.
- P circulates a COORDINATOR message with the new high number.
- Here, both 2 and 1 elect 6:

```
graph BT; In1[ ] --> 0((0)); In2[ ] --> 0; 0 -- "[5,6,0]" --> 1((1));
```

[5, 6, 0, 1, 2, 3, 4]  
[2, 3, 4, 5, 6, 0, 1]





# Complexity

---

- Given we have  **$n$**  processes
- The complexity is always  **$2n$**